

Spooky Author Identification

Chintan Shah
Shah.c@husky.neu.edu

Introduction:

In this year's Halloween playground Kaggle competition, you're challenged to predict the author of excerpts from horror stories by Edgar Allan Poe, Mary Shelley, and HP Lovecraft. Pages from unbound spooky books are scattered across the chamber floor. The task is to figure out how to put them back together according to the authors who wrote them.

DataSet:

Dataset is provided by the Kaggle @ <https://www.kaggle.com/c/spooky-author-identification/data>
It consists of 2 set of data as below :

1. Train.csv : This set contains 19579 rows & 3 columns which are : id, text & authors related to the text.

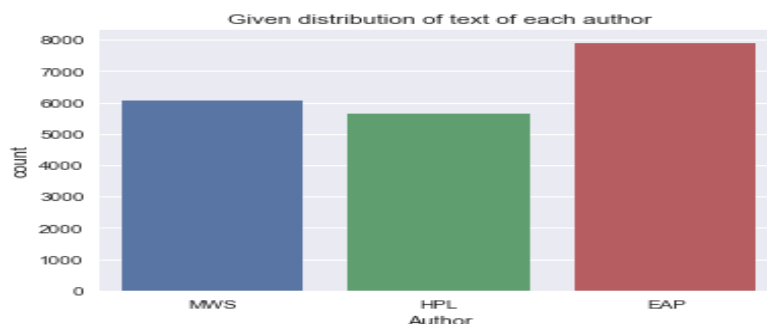
| id | text | author |
|---------|---|--------|
| id00001 | Idris was well content with this resolve of mine. | MWS |
| id00002 | I was faint, even fainter than the hateful mod... | HPL |
| id00003 | Above all, I burn to know the incidents of you... | EAP |
| id00004 | He might see, perhaps, one or two points with ... | EAP |
| id00005 | All obeyed the Lord Protector of dying England... | MWS |

2. Test.csv : This set contains 8392 & 2 columns which are : id, text.
The task is to find the authors of these text.

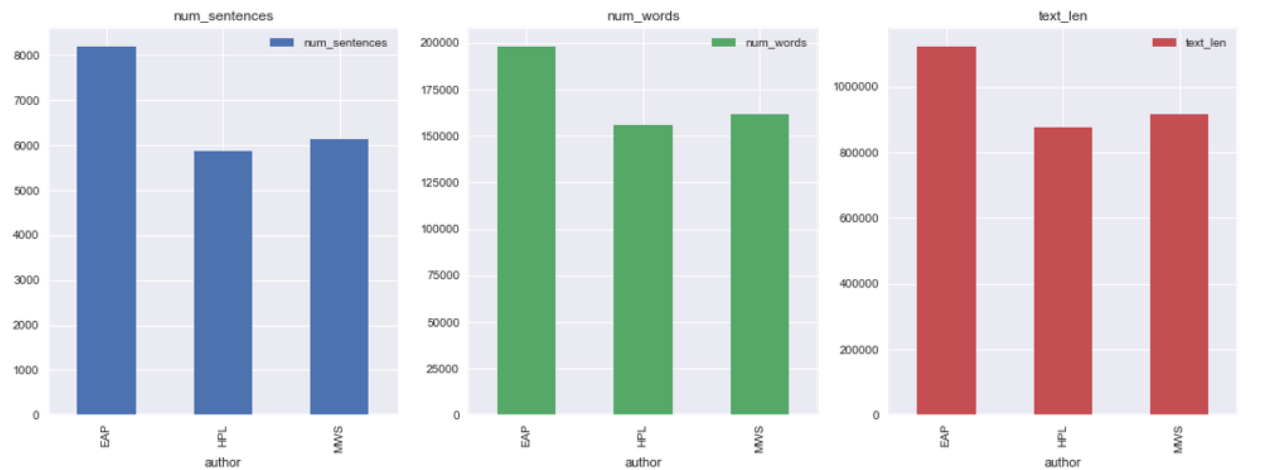
| id | Text |
|---------|---|
| id02310 | Still, as I urged our leaving Ireland with suc... |
| id24541 | If a fire wanted fanning, it could readily be ... |
| id00134 | And when they had broken down the frail door t... |
| id27757 | While I was thinking how I should possibly man... |
| id04081 | I am not sure to what limit his knowledge may ... |

Exploratory data analysis:

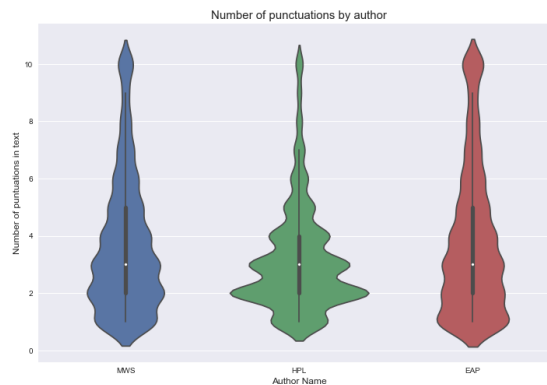
1. Distribution of text on basis of Authors : We have maximum text from EAP



2. Features of text used by each authors :



3. Punctuation Used by the authors :



4. Gram tables of each Author :

a. EAP :

| | 1-Gram | Occurence | 2-Gram | Occurence | 3-Gram | Occurence | 4-Gram | Occurence |
|---|--------|-----------|-------------|-----------|-------------------|-----------|-------------------------------|-----------|
| 0 | upon | 1025 | let u | 50 | madame l espanaye | 12 | general john b c | 12 |
| 1 | one | 671 | three four | 23 | general john b | 12 | brigadier general john b | 7 |
| 2 | could | 453 | l etoile | 23 | john b c | 12 | ugh ugh ugh ugh | 6 |
| 3 | would | 409 | one two | 23 | ha ha ha | 11 | brevet brigadier general john | 6 |
| 4 | said | 356 | every thing | 22 | barrière du roule | 10 | john b c smith | 6 |

b. HPL :

| | 1-Gram | Occurence | 2-Gram | Occurence | 3-Gram | Occurence | 4-Gram | Occurence |
|---|--------|-----------|-----------|-----------|--------------------|-----------|-----------------------------|-----------|
| 0 | one | 516 | old man | 59 | heh heh heh | 9 | eric moreland clapham lee | 4 |
| 1 | could | 480 | could see | 31 | terrible old man | 9 | oonai city lute dancing | 4 |
| 2 | thing | 433 | one night | 23 | charles le sorcier | 8 | aira city marble beryl | 4 |
| 3 | old | 392 | old woman | 22 | small paned window | 8 | mad arab abdul alhazred | 4 |
| 4 | would | 357 | one might | 19 | great great great | 6 | necronomicon mad arab abdul | 4 |

c. MWS:

| | 1-Gram | Occurence | 2-Gram | Occurence | 3-Gram | Occurence | 4-Gram | Occurence |
|---|--------|-----------|-----------------|-----------|-------------------|-----------|-------------------------|-----------|
| 0 | one | 489 | old man | 29 | let u go | 4 | next day next hour | 2 |
| 1 | would | 475 | lord raymond | 28 | whole human race | 3 | nearest town took post | 2 |
| 2 | could | 383 | fellow creature | 23 | became every day | 3 | abode near hyde park | 2 |
| 3 | life | 350 | one day | 21 | time lord raymond | 3 | five year old mother | 2 |
| 4 | yet | 318 | let u | 16 | might one day | 3 | like thousand pack wolf | 2 |

Pre-processing & Model Creation:

To predict the author, I have used the classification Algorithms and build the model to predict the probabilities of each author & have save the result in a csv file. So before creating the model, there is some pre-processing that was needed so as to fit the data into the model.

To start with the pre-processing of text I have used Natural Language Toolkit(NLTK).

1. Natural Language Toolkit :

NLTK is a leading platform for building Python programs to work with human language data. It provides easy-to-use interfaces to over 50 corpora and lexical resources such as WordNet, along with a suite of text processing libraries for classification, tokenization, stemming, tagging, parsing, and semantic reasoning, wrappers for industrial-strength NLP libraries.

Using NLTK following things are preformed :

- a. Remove Punctuation.
- b. Removing stopwords : Stop Words are words which do not contain important significance
- c. Stemming & Lemmatization : Lemmatization usually refers to doing things properly with the use of a vocabulary and morphological analysis of words, normally aiming to remove inflectional endings only and to return the base or dictionary form of a word, which is known as the lemma

After Process the text using the NLTK we need to convert the text in the form of sparse matrix so that it can be used by the algorithms.

2. CountVectorizer :

Convert a collection of text documents to a matrix of token counts

This implementation produces a sparse representation of the counts using `scipy.sparse.csr_matrix`.

Countvectorizer gives out the sparse matrix which contains the frequency of each word in respective documents/Data

But on the basis of the count of words, only the words that have appeared for most number of time is given the importance due to which we may miss out some important information that is present with low frequency words. So to balance out this, I have also used TF-IDF

3. Term Frequency–Inverse Document Frequency (TF-IDF) :

TF: Term Frequency, which measures how frequently a term occurs in a document. Since every document is different in length, it is possible that a term would appear much more times in long documents than shorter ones. Thus, the term frequency is often divided by the document length (aka. the total number of terms in the document) as a way of normalization:

$$TF(t) = (\text{Number of times term } t \text{ appears in a document}) / (\text{Total number of terms in the document}).$$

IDF: Inverse Document Frequency, which measures how important a term is. While computing TF, all terms are considered equally important. However it is known that certain terms, such as "is", "of", and "that", may appear a lot of times but have little importance. Thus we need to weigh down the frequent terms while scale up the rare ones, by computing the following:

$$IDF(t) = \log_e(\text{Total number of documents} / \text{Number of documents with term } t \text{ in it}).$$

4. Applying the pre-process data on the classifier algorithms & measure the accuracy :

Below is the table with accuracy of different algorithms :

| Classification Algorithms | Accuracy using CountVectorizer | Accuracy using TF-IDF |
|---------------------------|--------------------------------|-----------------------|
| Logistic Regression | 0.789581205 | 0.801583248 |
| Navie Bayes | 0.822778345 | 0.813329928 |
| Simple SVM | 0.665219612 | 0.701736466 |
| K nearest neighbors | 0.618009091 | 0.593462717 |
| decision-tree algorithm | 0.593462717 | 0.593462717 |
| Random Forest | 0.665730337 | 0.683350358 |

As we see Navie Bayes has the best result with both Count Vectorizer& TF-IDF.

Deep Learning :


In the era of Deep Learning, I have also tried to predict the author using the Neural network concept and have used 2 models: A simple Neural Network & A Simple dense LSTM with GloVe vector representations for words.

a. Simple 4 Layer Network :

I have used the simple NN layers as below :

| Layer (type) | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_1 (Dense) | (None, 300) | 90300 |
| dropout_1 (Dropout) | (None, 300) | 0 |
| dense_2 (Dense) | (None, 300) | 90300 |
| dropout_2 (Dropout) | (None, 300) | 0 |
| dense_3 (Dense) | (None, 300) | 90300 |
| dropout_3 (Dropout) | (None, 300) | 0 |
| dense_4 (Dense) | (None, 3) | 903 |
| activation_1 (Activation) | (None, 3) | 0 |
| Total params: 271,803 | | |
| Trainable params: 271,803 | | |
| Non-trainable params: 0 | | |

Using the simple NN we got the below accuracy & loss :



Using the simple NN model we didn't had a great result & so I decided to use the LSTM NN.

b. Long short-term memory (LSTM) :


I have used the embedding layer with LSTM layer & then fully connected layers.

Below is the model summary :

| Layer (type) | Output Shape | Param # |
|---------------------------------------|-----------------|---------|
| embedding_1 (Embedding) | (None, 70, 300) | 7783200 |
| spatial_dropout1d_1 (Spatial Dropout) | (None, 70, 300) | 0 |
| lstm_1 (LSTM) | (None, 100) | 160400 |
| dense_5 (Dense) | (None, 1024) | 103424 |
| dropout_4 (Dropout) | (None, 1024) | 0 |
| dense_6 (Dense) | (None, 1024) | 1049600 |
| dropout_5 (Dropout) | (None, 1024) | 0 |
| dense_7 (Dense) | (None, 3) | 3075 |
| activation_2 (Activation) | (None, 3) | 0 |
| Total params: 9,099,699 | | |
| Trainable params: 1,316,499 | | |
| Non-trainable params: 7,783,200 | | |

Using the LSTM model I have got much better result than the previous NN model.

Accuracy & loss v/s Epochs :



As we see after each epoch the accuracy increases and also the loss decreases

Conclusion:

So with Classifiers we have the best result when we used Naïve Bayes Classifier & also we see that when we used LSTM RNN model with 25epochs the accuracy increases & the loss decreases. In future, if more parameters are tuned in we could fetch much better result with Naïve bayes & LSTM models. Also in future we cn use XGBoost & GridSearch CV to produce better results.

References :

1. <https://www.youtube.com/watch?v=Tp3SaRbql4k> – different keras Layers.
2. <http://ruder.io/optimizing-gradient-descent/> - Gradient descent optimization algorithms
3. <https://www.youtube.com/watch?v=4rG8lsKdC3U> – how to implement LSTM
4. <https://www.youtube.com/watch?v=WCUNPb-5EYI> – A very good example which helps in understanding LSTM (example starts from 18.00)
5. <https://www.kaggle.com/abhishek/approaching-almost-any-nlp-problem-on-kaggle> - Approaching NLP problem.
6. <https://nlp.stanford.edu/projects/glove/> - GloVe Word Vector
7. <https://medium.com/mlreview/understanding-istm-and-its-diagrams-37e2f46f1714> - Article helps to understand LSTM.